AI TRANSFORMATION PROGRAM

Day 1

Strategy & Leadership Alignment

Facilitator & Participant Guide

End-to-end notes · diagrams · flowcharts · hands-on labs · use cases

Audience: Executives, business & technology leaders, transformation sponsors

Format: Instructor-led workshop — approx. 1 day (5 modules)

Pre-requisites: None technical; bring a real business challenge to work on

Version: June 2026 · current to Microsoft Build 2026 announcements

Contents

How to use this guide

This guide is written for both the trainer and the participant. The main text is the shared teaching content. Coloured boxes signal who a block is for and what to do with it:

FACILITATOR NOTE

For the trainer. Timings, talking points, facilitation moves, and what “good” looks like. Participants can read these too — there are no secrets — but they are written to help whoever is leading the room.

HANDS-ON LAB

Hands-on lab / activity. A structured exercise with steps, a worksheet, and a time-box. Every module has at least one. Blank worksheets are collected in Appendix B.

USE CASE

Use case. A concrete, often industry-specific example that grounds the concept in something real.

KEY TAKEAWAYS

- The three-to-five things a participant must remember from the section.

WATCH OUT

A common trap, anti-pattern, or risk to call out explicitly.

Master run sheet — the full 8-hour day

This guide is built for a full-day delivery: **approximately 8 hours of facilitated content** (480 minutes), wrapped in a 45-minute lunch and two 15-minute breaks for a 9:00–18:00 clock day. Clock times below assume a 9:00 start — shift them to fit your room. Each module has its own minute-by-minute run sheet in the next section.

Clock	Block	Min	Focus / outcome
09:00–09:30	**Welcome & framing**	30	Logistics, introductions, program arc, golden-thread setup
09:30–11:00	**Module 1 · AI Transformation Overview & Trends**	90	Shared language; an honest read of where the org sits
11:00–11:15	Break	15	—

Clock	Block	Min	Focus / outcome
11:15–12:45	**Module 2 · Use-Case Identification Workshop**	90	A prioritized shortlist of candidate use cases
12:45–13:30	Lunch	45	—
13:30–15:00	**Module 3 · Microsoft AI Ecosystem**	90	Map use cases to the right Microsoft building blocks
15:00–15:15	Break	15	—
15:15–16:35	**Module 4 · ROI, Cost Models & Value Realization**	80	A defensible business case for one use case
16:35–17:55	**Module 5 · Governance & Responsible AI**	80	A risk-tiering and governance approach leaders can adopt
17:55–18:15	**Wrap-up & action planning**	20	Synthesis, personal commitments, bridge to Day 2

Facilitated content: 30 + 90 + 90 + 90 + 80 + 80 + 20 = **480 minutes (8 hours)**. Breaks + lunch: 75 minutes. Total room time: 9 hours 15 minutes.

FACILITATOR NOTE

Running order. Modules 1–3 run before lunch on purpose — they build momentum and produce the use-case shortlist that Modules 4–5 stress-test. If you lose time, protect the labs (that is where learning sticks) and compress the teach blocks, not the other way round.

Compressing to a shorter day. For a 6-hour version, cut Module 1 to 60 (drop the trends deep-dive), run a single shared use case instead of per-table, and shorten Modules 4–5 labs to 10 minutes each. For a half-day executive briefing, deliver Modules 1, 4 and 5 only (teach, no labs).

Materials. Sticky notes, markers, a 2×2 grid on a wall or whiteboard, printed worksheets (Appendix B), one laptop per table with internet for the Module 3 lab, a projector for the diagrams, and a visible “parking lot” sheet for off-topic but valuable points.

The golden thread. Ask each table to pick ONE real business problem in Module 2 and carry it through every later module: map it to Microsoft tooling (M3), build its business case (M4), and tier its risk (M5). This single continuous example is what makes the day stick — reference it by name in every module.

What participants will be able to do by end of Day 1

1. Describe what AI transformation means for their organization and place it honestly on a maturity curve.
2. Distinguish predictive, generative, and agentic AI — and choose the right pattern (or none) for a problem.
3. Identify, qualify, and prioritize business use cases using a repeatable value×feasibility method.
4. Map a use case to the Microsoft AI stack and decide whether to buy, configure, or build.
5. Construct a credible cost model and ROI narrative, and define metrics that prove value.

6. Apply Responsible-AI principles and a risk-tiered governance model with clear decision gates.

Detailed facilitator run sheets

Each block below is broken into timed segments so you can run the day to the clock. Times are *minutes into that module* (00:00 = the moment the block starts). The bold **Talk track** boxes are sample words for the highest-leverage moments — adapt them to your voice; do not read them out.

Welcome & framing · 30 minutes (09:00–09:30)

Time	Segment	Min	What the facilitator does
00:00–10:00	Welcome, logistics, ground rules	10	Housekeeping (exits, Wi-Fi, timing). Set ground rules: phones on silent, one conversation at a time, “yes-and” not “no-but”, the parking lot. Name the goal of the day in one line.
10:00–20:00	Introductions	10	Round the room (or table-by-table if >20 people): name, role, and one sentence — “the one thing about AI that excites you and the one thing that worries you.” Capture themes on a flip chart; you will call back to them all day.
20:00–30:00	Program arc & golden-thread setup	10	Show the master run sheet. Explain the five modules build on each other. Plant the golden thread: “By the first break you will each pick ONE real problem and we will carry it all the way to a costed, risk-tiered decision by end of day.”

TALK TRACK — SAY SOMETHING LIKE THIS

“Before we talk about technology, a confession: most AI programs fail not because the models are weak, but because they get pointed at the wrong problems with no business case and no guardrails. Today is about fixing that order of operations. By 6 o’clock each of you will leave with one real use case from your own world — qualified, mapped to Microsoft tooling, costed, and risk-tiered. Same example, five lenses. Let’s begin with where the industry actually is — not the hype.”

Module 1 run sheet · 90 minutes

Time	Segment	Min	What the facilitator does
00:00–10:00	Opening hook + objectives	10	Run the “raise your hand if you use AI / keep it up if AI has changed a core process” hook. Most hands drop — that gap is the whole module. State the four objectives.
10:00–25:00	Adoption vs transformation	15	Teach §1.1 with the comparison table. Ask: “Which row is your org honestly in today?” Take 3–4 answers.

Time	Segment	Min	What the facilitator does
25:00–40:00	Three waves of AI	15	Walk Figure 1.1. For each wave ask the room for one example they already use. Stress the waves coexist — pick the simplest that solves the problem.
40:00–1:05	Maturity curve + self-assessment	25	Show Figure 1.2, then run the maturity self-assessment lab. Tables score themselves and name the single next concrete step. Debrief 3 tables.
1:05–1:20	2026 trends, read critically + winners	15	Teach §1.4–1.5. Emphasize the adoption-vs-production gap and that winners redesign workflows, not just buy licenses. Cite figures as ranges.
1:20–1:30	Knowledge check + takeaways	10	Run the knowledge check verbally or per-table; land the key takeaways. Bridge to Module 2.

Module 2 run sheet · 90 minutes

Time	Segment	Min	What the facilitator does
00:00–10:00	Why most use-case lists fail	10	Open with the failure pattern (vendor-led wish-lists, no qualification). Set the objective: leave with a prioritized shortlist.
10:00–22:00	Five qualification lenses	12	Teach the five lenses. Have each table hold up their golden-thread problem against each lens as you go.
22:00–32:00	Pattern fit & disqualifiers	10	Walk the qualification flowchart (Figure 2.1). Make the disqualifiers explicit — most weak ideas die here, which is the point.
32:00–42:00	Priority matrix method	10	Teach value×feasibility (Figure 2.2). Demonstrate scoring one example live on the wall grid.
42:00–1:25	Live workshop	43	Tables generate 5–8 candidates, qualify, plot on the 2×2, pick ONE, and fill the one-page brief. Circulate; push on vague value statements.
1:25–1:30	Debrief + takeaways	5	Two tables share their chosen use case and why. Confirm everyone has a written brief — this is the golden thread for the rest of the day.

TALK TRACK — SAY SOMETHING LIKE THIS

Launching the workshop: “Here is the discipline: an idea is not a use case until it survives five questions. We are not looking for the most exciting idea — we are looking for the one with real

> value that you can actually feasibly do. In 45 minutes I want one sentence from each table: ‘We will help [who] do [what] so that [measurable outcome].’ If you can’t say the measurable outcome, it is not ready. Pick the boring, winnable one over the flashy, vague one.”

Module 3 run sheet · 90 minutes

Time	Segment	Min	What the facilitator does
00:00–12:00	Currency caveat + stack overview	12	Flag that names/pricing change fast (verify on Microsoft Learn). Walk Figure 3.1 top to bottom — the five layers wrapped in governance.
12:00–24:00	Microsoft Foundry	12	Cover the unified build platform: model catalog, agent service, evaluations, retrieval/Foundry IQ. Keep it conceptual — leaders need the shape, not SDK syntax.
24:00–39:00	Copilot family + Agent 365	15	Walk the Copilot table: M365 Copilot, Copilot Studio, agents, and Agent 365 as the governance plane. Anchor each to “buy vs configure vs build.”
39:00–49:00	Fabric & Fabric IQ	10	Explain OneLake, Copilot in Fabric, and Fabric IQ as the semantic/ontology layer that gives agents shared business meaning.
49:00–57:00	Buy / configure / build decision	8	Teach the decision table. Map the worked example onto it live.
57:00–1:27	Stack-mapping lab	30	Each table maps its golden-thread use case onto the stack: which layer, which product, buy/configure/build, and what data it needs. Circulate.
1:27–1:30	Check + takeaways	3	Quick knowledge check; land the takeaways. Bridge to ROI.

Module 4 run sheet · 80 minutes

Time	Segment	Min	What the facilitator does
00:00–08:00	Why ROI is hard for AI	8	Name the traps: soft benefits, value that leaks between pilot and production, variable token costs. Set up that we price the whole stack.
08:00–20:00	The full cost stack	12	Teach the cost table — build, run, change management, and variable usage. Stress the costs people forget (redeployment, governance, model spend at scale).
20:00–30:00	Benefit categories	10	Cover cost-avoidance, revenue, risk reduction, and quality. Every benefit

Time	Segment	Min	What the facilitator does
			needs a category, a number, and a baseline.
30:00–38:00	Value-realization loop	8	Walk Figure 4.1 (measure → learn → reinvest). Value is realized over time, not at launch.
38:00–53:00	Worked business case	15	Walk the worked example end to end (≈$208K net). Pause at each assumption and ask “would your CFO accept this?”
53:00–1:12	Build-your-case lab	19	Tables build a rough cost–benefit and one-line value narrative for their golden-thread use case using worksheet B5.
1:12–1:20	Metrics + check + takeaways	8	Define the leading/lagging metrics that prove value; run the check; land takeaways.

TALK TRACK — SAY SOMETHING LIKE THIS

Framing the worked case: “Watch where the money actually moves. The license cost is the easy part — the parts that sink business cases are the ones nobody put in the spreadsheet: change management, the variable usage cost at ten-times adoption, and whether the time you ‘saved’ was actually redeployed to something valuable. A model that survives the question ‘where did the saved time go?’ is fundable. One that doesn’t, isn’t.”

Module 5 run sheet · 80 minutes

Time	Segment	Min	What the facilitator does
00:00–08:00	Why governance now	8	Make the case that governance is an enabler, not a brake — it is what lets you ship agents safely and at speed. Note the rise of agent-linked incidents.
08:00–20:00	Six Responsible-AI principles	12	Walk Figure 5.1. For each principle give one concrete “what it means on Monday” example.
20:00–32:00	Operating model + decision gates	12	Teach Figure 5.2: the four-tier org and the G1–G4 gates. Show where each gate sits in a project’s life.
32:00–42:00	Risk tiering	10	Teach low/medium/high tiering and how the tier sets the depth of review. Tier the worked example together.
42:00–50:00	Microsoft tooling + agentic risks	8	Map Entra, Purview, Defender, Agent 365 to the principles. Call out agent-specific risks (autonomy, tool access, identity).
50:00–1:08	Risk-tier lab + scenarios	18	Tables tier their golden-thread use case, name an accountable owner, and pick the

Time	Segment	Min	What the facilitator does
			gates it must pass (worksheet B6). Run 1–2 scenario cards.
1:08–1:20	Check + takeaways	12	Run the check; land takeaways; set up the wrap-up synthesis.

Wrap-up & action planning · 20 minutes (17:55–18:15)

Time	Segment	Min	What the facilitator does
00:00–08:00	Golden-thread synthesis	8	Pick one table and walk their use case through all five lenses out loud: problem → pattern → stack → business case → risk tier. This is the “aha” that proves the day connected.
08:00–15:00	Personal action items	7	Each person writes two commitments: one thing to validate and one conversation to have within two weeks. Have a few read theirs aloud.
15:00–20:00	Bridge to Day 2 + close	5	Preview Day 2 (deeper build / pilots). Thank the room; collect worksheets if you want the artifacts.

Module 1 · AI Transformation Overview and Industry Trends

Duration: 90 minutes · **Format:** Teach + facilitated discussion + self-assessment

Learning objectives

1. Separate "AI adoption" (using tools) from "AI transformation" (redesigning how work happens).
2. Explain the three waves — predictive, generative, agentic — and why each demands more governance.
3. Read the 2026 industry landscape critically, including the gap between adoption and production value.
4. Locate your organization on the AI maturity curve and name the next concrete step.

FACILITATOR NOTE

90-minute plan. (1) Hook & definitions – 15 min. (2) Three waves – 15 min. (3) Maturity curve + self-assessment activity – 25 min. (4) 2026 trends – 20 min. (5) "Winners vs stallers" discussion – 10 min. (6) Knowledge check & bridge – 5 min.

Opening hook. Ask the room: "Raise your hand if your organization is using AI." (Most hands go up.) Then: "Keep it up if AI has changed a core process — not just sped up an email." (Most hands drop.) That gap is the entire reason for this program.

1.1 Adoption is not transformation

Most organizations have *adopted* AI: people use a chat assistant, a copilot drafts a slide, a model scores a lead. That is real and useful, but it is incremental — the underlying process is unchanged. *Transformation* is when AI changes the shape of the work itself: steps disappear, decisions move closer to the data, cycle times collapse, and new capabilities become possible that simply could not exist before.

The distinction matters because the two are governed, funded, and measured differently. Adoption is a productivity line item. Transformation is a strategic bet that touches org design, data, risk, and culture — which is exactly why Day 1 is aimed at leadership, not just builders.

Dimension	AI adoption	AI transformation
Unit of change	A task or a person	A process or a business model
Sponsor	Team lead / individual	Executive / board
Measure of success	Time saved per user	Outcome moved (cost, revenue, risk, CX)
Main risk	Under-use, sprawl	Mis-prioritization, governance debt
Typical 2026 reality	Widespread	Concentrated in a few leaders

1.2 Three waves of enterprise AI

Three capability waves now coexist. They are additive, not sequential replacements — a mature program runs all three, choosing the simplest pattern that solves the problem.

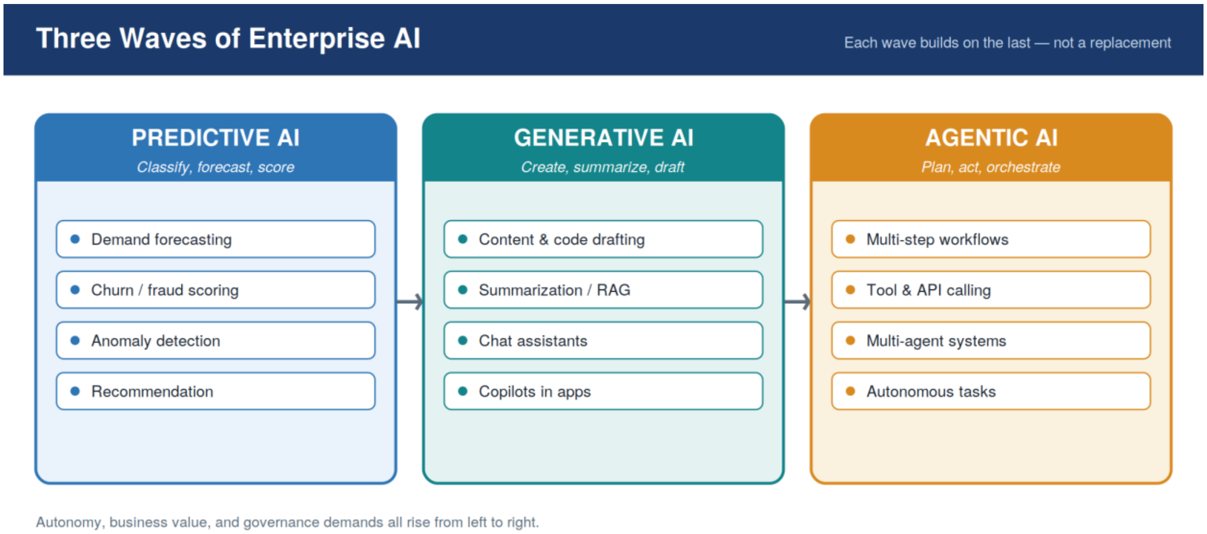

Figure 1.1 — Three waves of enterprise AI. Autonomy, value, and governance demands all rise left to right.

- **Predictive AI** classifies, forecasts, and scores. Mature, well-understood, often the highest-ROI place to start (fraud, churn, demand, anomaly detection).
- **Generative AI** creates and summarizes — drafting, RAG over enterprise knowledge, assistants and copilots embedded in the tools people already use.
- **Agentic AI** plans and acts across multiple steps, calling tools and APIs, sometimes coordinating multiple agents. This is the defining shift of 2025–26 — and the one with the steepest governance curve.

WATCH OUT

“Agentic” is not automatically better. Industry guidance is explicit: use agents where multi-step autonomy genuinely adds value, automation for routine workflows, and a simple assistant for retrieval. Reaching for an autonomous agent when a rule or a single model call would do is a fast route to cost, fragility, and risk.

1.3 The AI transformation maturity curve

Transformation is a journey through repeatable stages. Naming the stage you are in is the single most useful thing leadership can do on Day 1, because each stage has a different “next move.”

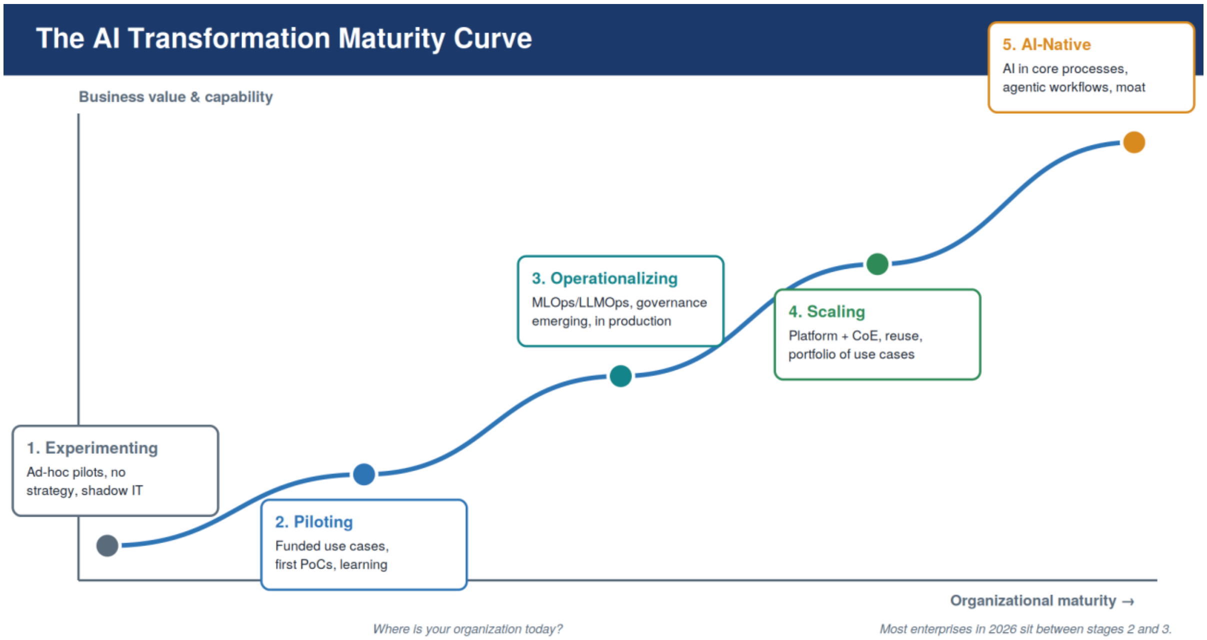

Figure 1.2 — The AI transformation maturity curve.

Stage	You'll recognize it by…	The next move
1 · Experimenting	Scattered pilots, shadow tools, no shared strategy	Pick a sponsor and 2–3 funded use cases
2 · Piloting	Real PoCs, early wins, but nothing in production	Get one use case into production with monitoring
3 · Operationalizing	Things run in production; governance is emerging	Stand up a platform + lightweight governance gates
4 · Scaling	A CoE, reuse, a portfolio, measured value	Industrialize: reusable patterns, FinOps, talent
5 · AI-native	AI in core processes, agentic workflows, a moat	Protect the moat; push into new business models

In 2026 most enterprises sit between stages 2 and 3 — plenty of pilots, far less in production. Closing that gap is where value is won or lost (see 1.4).

1.4 Industry trends 2026 — read critically

The headline numbers are dramatic; the useful signal is in the gaps. A facilitator should present these as orientation, not gospel — figures move every quarter and sources disagree.

Agents have gone mainstream — in shipping software, not always in production

- **Embedding is now the default.** Analysts estimate a large share of enterprise applications shipped or updated in early 2026 embed at least one task-specific agent — up sharply from a couple of years earlier (Gartner).
- **But the production gap is the story of 2026.** Across surveys, a large majority of enterprises have adopted agents in some form while only a minority run them in production. That backlog — not the hype — is where competitive advantage is decided (IDC / McKinsey / S&P, ranges vary).

ROI is real but uneven

- Returns can be strong: IDC/Microsoft have cited an average return of roughly 3.7x per $1 invested in generative AI; Forrester analysis of Microsoft 365 Copilot has reported triple-digit three-year ROI with several hours saved per user per month.
- …but most initiatives disappoint: an IBM CEO study found only about a quarter of AI initiatives delivered the ROI expected, and Gartner has warned that a large share of agentic-AI projects may be cancelled by 2027 for unclear value, cost, or risk.
- What separates the two: McKinsey's high performers — the small group attributing a meaningful share of EBIT to AI — invest a larger fraction of digital budgets in AI and, critically, **redesign workflows around it rather than bolting AI onto old ones.**

New risk surface

- Agent autonomy, broad data access, tool-calling, and memory create an attack and failure surface most security teams are not yet equipped for. Analysts have begun tracking agent-linked breaches and "harm" incidents, and regulators are paying attention.

WATCH OUT

Facilitator caution on statistics. Quote ranges, not single hero numbers, and attribute the source and roughly the date. Half of these figures will have shifted by next quarter. The durable insight — "widespread adoption, scarce production value, value comes from workflow redesign and governance" — will not.

1.5 Why some organizations pull ahead

Strip away the tooling and the leaders share a recognizable operating model:

1. They start from business problems, not from models — a prioritized portfolio, not a pile of pilots.
2. They redesign the workflow, not just insert AI into the existing one.
3. They treat data as the prerequisite — governed, accessible, semantically described.
4. They build a thin platform and a CoE so the second and tenth use cases are cheaper than the first.
5. They govern from day one — light enough not to block, real enough to be trusted.
6. They measure outcomes, kill what doesn't work, and reinvest in what does.

Use cases — cross-industry snapshot

USE CASE

How the same three waves show up across industries:

- **Banking & insurance (front-runners):** predictive fraud/risk scoring → generative claim and KYC summarization → agentic onboarding and dispute handling.
- **Manufacturing:** predictive maintenance → generative SOP and engineering-doc assistants → agentic supply-chain replanning.
- **Healthcare & public sector (trailing, high-stakes):** predictive triage/no-show models → generative documentation → cautious, heavily-governed agentic pilots.

- **Retail:** demand forecasting → generative product content and support → agentic merchandising and service resolution.

Hands-on · Maturity self-assessment

HANDS-ON LAB

Goal. Place your organization on the maturity curve honestly, and name one concrete next move. 20 minutes.

Steps.

1. Individually, score your organization 1–5 on each of: strategy, data readiness, people/skills, governance, production footprint (worksheet B1).
2. Take the lowest of your five scores as your honest stage — you are only as mature as your weakest pillar.
3. At your table, compare scores. Where you disagree by 2+ points, discuss why (often a sign that “AI” means different things to different functions).
4. Agree one sentence: “We are at stage __, and our single next move is ____.”

Debrief. Each table reads its sentence aloud. Capture them — they become the reality check for the rest of the day.

FACILITATOR NOTE

Answer guidance. There is no “correct” stage — the value is in the honesty. Two facilitation moves: (a) gently challenge any table that rates itself stage 4–5 but can’t name a production system with monitoring; (b) reassure stage 1–2 tables that this is the norm — the goal is the next step, not a leap to AI-native.

Knowledge check

1. In one sentence, what is the difference between AI adoption and AI transformation?
2. Give one example each of a predictive, a generative, and an agentic use case from your own organization.
3. Which maturity stage are you in, and what is the single biggest blocker to the next stage?
4. Why is “more agents” not automatically better? Give a case where a simpler pattern is the right answer.

KEY TAKEAWAYS

- Adoption ≠ transformation. Transformation redesigns the work; it is a leadership bet.
- Predictive, generative, and agentic AI coexist — pick the simplest pattern that solves the problem.
- Name your maturity stage honestly; you are only as mature as your weakest pillar.
- 2026’s real story is the adoption-to-production gap. Value comes from workflow redesign + governance, not tools.

Module 2 · Business Use-Case Identification Workshop

Duration: 90 minutes · **Format:** Short teach blocks wrapped around a live, hands-on workshop

Learning objectives

1. Generate candidate use cases systematically from the value chain, not randomly.
2. Qualify a candidate with a repeatable checklist — including when AI is the wrong answer.
3. Prioritize using a value × feasibility matrix the whole leadership team can agree on.
4. Capture a use case on a one-page brief that survives the journey to a business case.

FACILITATOR NOTE

This module is mostly workshop. Keep teaching to ~30 minutes total and protect ~50 minutes for the activity. The deliverable — a prioritized shortlist per table — feeds Modules 3, 4 and 5, so do not let it slip.

Timing. (1) Where good use cases come from – 10 min. (2) Pattern fit & disqualifiers – 8 min. (3) The qualification flow – 7 min. (4) Workshop – 50 min (brainstorm 12 / qualify 10 / score & plot 15 / brief & pitch 13). (5) Debrief + knowledge check – 15 min.

2.1 Where good use cases come from

Strong candidates are not brainstormed in the abstract — they are found by walking the business systematically and looking for friction that AI is genuinely suited to remove.

Lens	Question to ask	Signals a good candidate
Value chain	Which steps are slow, costly, or error-prone?	High volume + repetitive + measurable
Pain points	Where do people complain or wait?	Manual rekeying, long queues, backlogs
Data-rich	Where does data already pile up unused?	Logs, tickets, documents, transactions
Decisions	Which judgements are made at scale on patterns?	Scoring, routing, triage, forecasting
Knowledge	Where is expertise scarce or hard to find?	“Ask the one expert” bottlenecks

2.2 Pattern fit — and when not to use AI

Match the problem to the simplest pattern. Then sanity-check that AI is warranted at all.

- **Use predictive** when the job is to forecast or score from historical patterns.
- **Use generative** when the job is to draft, summarize, translate, or answer from a knowledge base.
- **Use agentic** when the job is genuinely multi-step, requires calling tools/systems, and benefits from autonomy.

WATCH OUT

Disqualifiers — reach for rules, automation, or BI instead when:

- The logic is deterministic and well-specified (a rule engine is cheaper, faster, auditable).
- A wrong answer is unacceptable and there is no human in the loop (high stakes, no oversight).
- The data needed doesn’t exist, isn’t accessible, or is too poor to trust.
- The value is trivial — you’ll spend more governing it than it returns.

2.3 Qualification flowchart

Run each candidate through the same gates. Anything that survives is worth scoring; anything that fails tells you what to fix first (data, controls, or the business case).

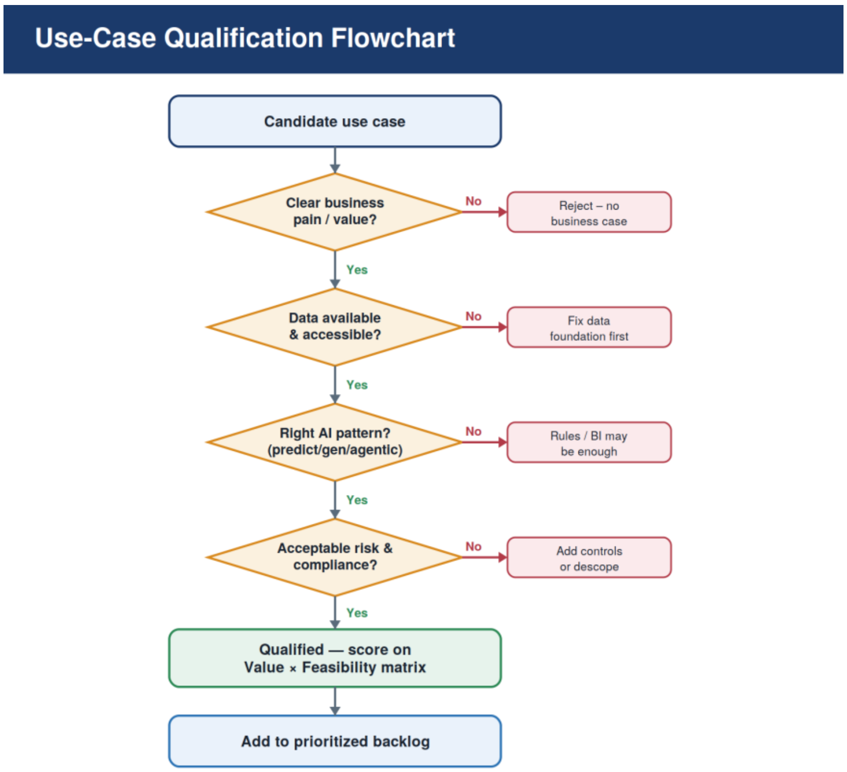

Figure 2.1 — Use-case qualification flowchart.

2.4 Prioritize: value × feasibility

Qualified candidates are plotted on two axes the leadership team scores together. The conversation the matrix forces is as valuable as the placement itself.

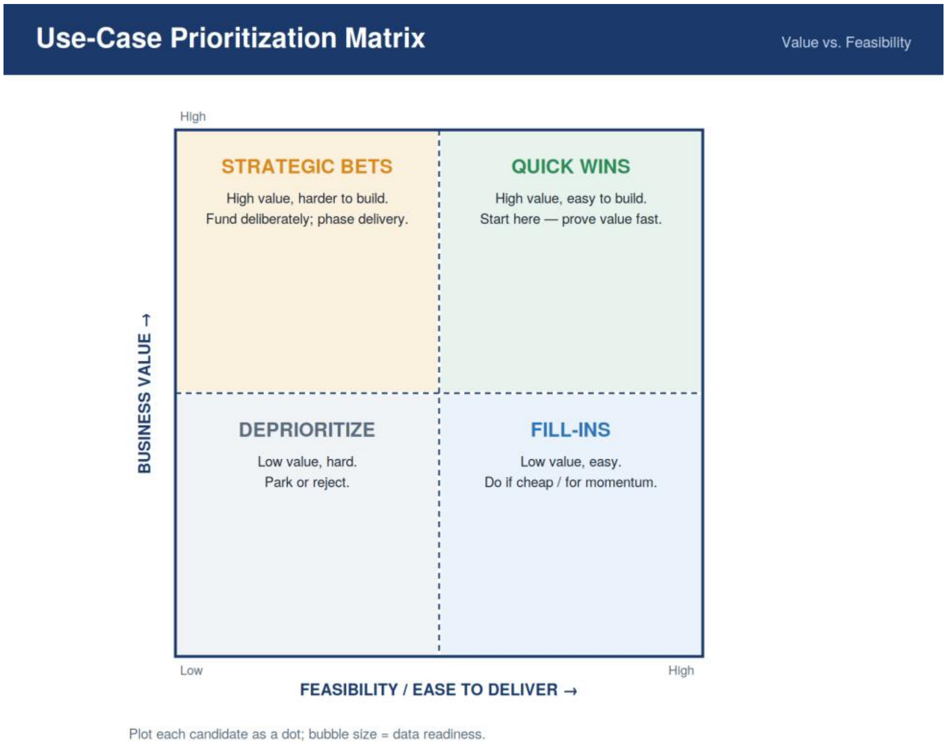

Figure 2.2 — Use-case prioritization matrix.

- **Quick wins** (high value, easy): start here to build credibility and learning.
- **Strategic bets** (high value, hard): fund deliberately and phase — these are where transformation lives.
- **Fill-ins** (low value, easy): do only if nearly free or useful for momentum/skills.
- **Deprioritize** (low value, hard): park or reject without guilt.

Score value and feasibility 1–5 each. Suggested factors — **value:** size of benefit, strategic fit, number of people affected; **feasibility:** data readiness, technical complexity, change-management load, risk/compliance burden.

2.5 The one-page use-case brief

Every shortlisted use case gets a one-pager. It is the artefact that carries forward into Module 4 (business case) and Module 5 (risk tiering).

Field	What goes here
Title & owner	Short name; the accountable business owner (not IT)
Problem	The business pain, in one or two sentences, with a number if possible
Today's process	How it works now; where the friction is
AI pattern	Predictive / generative / agentic — and why
Data needed	Sources, owner, accessibility, quality, sensitivity
Success metric	The ONE number that must move, and its baseline
Value & feasibility	Scores (1–5 each) and the quadrant

Field	What goes here
Risks / compliance	Top risks; any regulated data or decisions

Hands-on · The use-case workshop (core activity)

HANDS-ON LAB

Goal. Produce a prioritized shortlist of 2–3 qualified use cases per table, each on a one-page brief. ~50 minutes.

You will need. Sticky notes, markers, the printed brief (B2) and scorecard (B3), and a 2×2 grid drawn on a flip-chart or wall.

Step 1 · Diverge (12 min). Using the five lenses in 2.1, each person silently writes one use case per sticky note — aim for 6–8 per person. No debate yet. Stick them all on the wall.

Step 2 · Cluster & qualify (10 min). Group duplicates. Run the top ~8 through the qualification flow (Fig 2.1). Discard anything that fails a gate; note why.

Step 3 · Score & plot (15 min). For each survivor, agree a value score and a feasibility score (1–5) using the scorecard, then place it on the 2×2. Bubble size = data readiness.

Step 4 · Brief & pitch (13 min). Pick your top 2–3 (favour quick wins + one strategic bet). Fill a one-page brief for each. Prepare a 60-second pitch.

FACILITATOR NOTE

Facilitation moves.

- Enforce silent divergence in Step 1 — it prevents the loudest voice from anchoring the room.
- In scoring, when a table can't agree, ask "what would change your mind?" — the disagreement usually reveals a missing fact (data, cost, or risk).
- Push back on "everything is high value, high feasibility." Force ranking: "if you could only do ONE next month, which?"
- Watch for solution-first sticky notes ("use a chatbot"). Redirect to the problem: "what pain does that remove, and for whom?"

USE CASE

Worked example — a B2B services firm.

Problem: support agents spend ~40% of their time searching policy and past tickets to answer customer emails; average first-response time is 9 hours.

Pattern: generative (RAG over the knowledge base + ticket history) — not agentic; a human still sends the reply.

Data: knowledge base + 3 years of resolved tickets, owned by Support Ops, reasonably clean, no special-category data.

Metric: first-response time (baseline 9h). Value 4, Feasibility 4 → quick win.

> **Why it survives the flow:** clear value, accessible data, right pattern, low risk (human-in-the-loop, no regulated decision). This is the example we carry into Modules 3–5.

Knowledge check

1. Name the five lenses for finding use cases. Which surfaced your best candidate?
2. Give one example where AI is the WRONG answer and a rule, automation, or BI report is better.
3. Why score value and feasibility separately rather than as a single “gut feel”?
4. What is the single success metric for your table’s top use case, and what is its baseline today?

KEY TAKEAWAYS

- Find use cases systematically (value chain, pain, data, decisions, knowledge) — not by brainstorming buzzwords.
- Qualify before you prioritize; the gates tell you what to fix (data, controls, business case).
- Value × feasibility forces the leadership conversation that actually allocates resources.
- Start with quick wins for credibility; fund one strategic bet for real transformation.

Module 3 · The Microsoft AI Ecosystem

Duration: 90 minutes · **Format:** Teach + tool tour + design lab

Note on currency. Microsoft's AI portfolio changes fast. This module reflects the landscape as of Microsoft Build 2026 (June 2026), including the rename of Azure AI Foundry to Microsoft Foundry. Always confirm current names, availability (GA vs preview), and pricing against Microsoft Learn before quoting them to a customer.

Learning objectives

1. Describe the Microsoft AI stack as five layers wrapped in governance.
2. Explain what Microsoft Foundry, the Copilot family, and Microsoft Fabric each do.
3. Decide, for a given use case, whether to buy a Copilot, configure in Copilot Studio, or build in Foundry.
4. Map your Module 2 use case onto the stack end-to-end.

FACILITATOR NOTE

Timing. (1) The stack overview – 10 min. (2) Foundry – 12 min. (3) Copilot family – 15 min. (4) Fabric & data – 10 min. (5) Buy/configure/build decision – 8 min. (6) Lab – 30 min. (7) Debrief + check – 5 min.

Keep it conceptual for a leadership audience. They don't need SDK syntax; they need to know which box solves which problem, what it costs roughly, and where governance plugs in. Resist going down rabbit holes — use the parking lot.

3.1 The stack at a glance

Think of Microsoft's AI offering as five layers, with a governance-and-security band running underneath all of them.

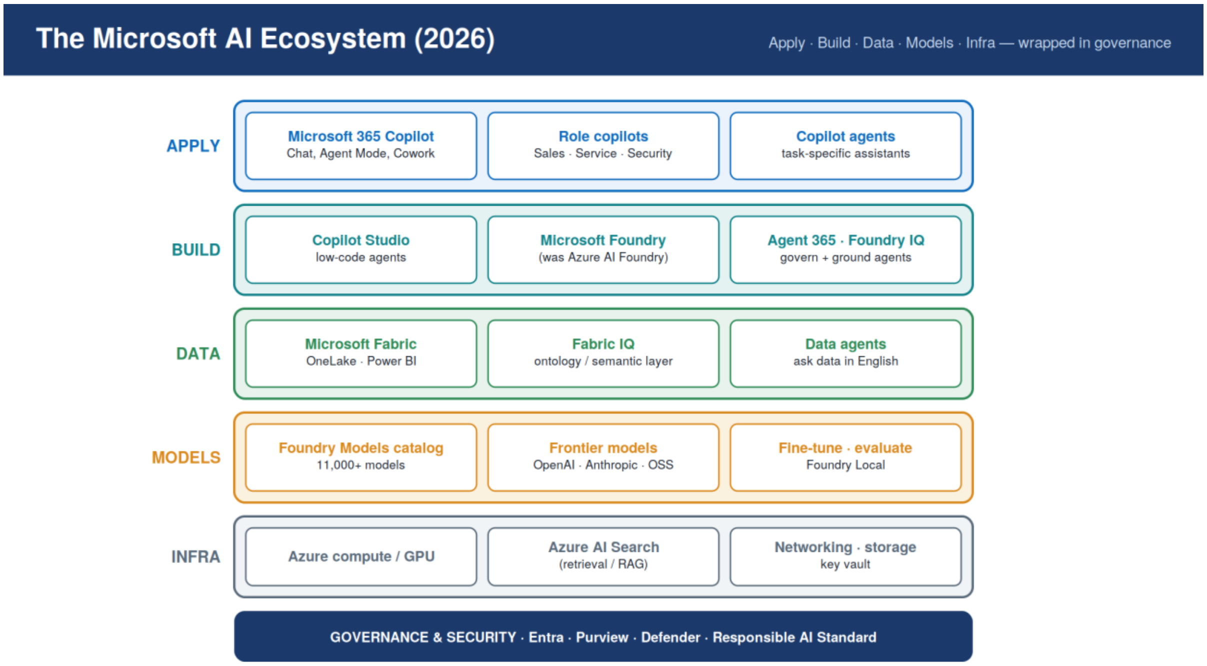

Figure 3.1 — The Microsoft AI ecosystem (2026).

- **Apply** — finished experiences employees use: Microsoft 365 Copilot (Chat, Agent Mode, Cowork), role copilots (Sales, Service, Security), and task-specific Copilot agents.
- **Build** — where you create custom agents and apps: Copilot Studio (low-code) and Microsoft Foundry (pro-code), governed by Agent 365 and grounded by Foundry IQ.
- **Data** — Microsoft Fabric and OneLake, with Fabric IQ as the semantic/ontology layer and data agents for natural-language analytics.
- **Models** — the Foundry model catalog (11,000+ models from Microsoft, OpenAI, Anthropic, Meta and others), plus fine-tuning, evaluation, and Foundry Local.
- **Infrastructure** — Azure compute/GPU, Azure AI Search for retrieval (RAG), networking, storage, and Key Vault.
- **Governance & security** — Microsoft Entra (identity), Microsoft Purview (data governance/compliance), Defender, and the Responsible AI Standard span every layer (Module 5).

3.2 Microsoft Foundry (formerly Azure AI Foundry)

Microsoft Foundry is the unified platform for building, evaluating, deploying and governing AI applications and agents. The platform was renamed from Azure AI Foundry effective January 2026 — the capabilities carried over unchanged.

- **Model catalog & routing** — discover, compare, fine-tune and deploy from 11,000+ models; route across models for cost/quality.
- **Foundry Agent Service** — build agents with tools, memory, and (as of 2026) multi-agent workflows, A2A and MCP connectivity, and managed network isolation.

- **Foundry IQ** — an SLA-backed retrieval/knowledge plane that grounds agents in enterprise data, reducing custom RAG plumbing.
- **Evaluations & observability** — built-in evaluation, tracing and monitoring; the place responsible-AI checks become testable (ties to Module 5).

Use Foundry when you need pro-code control, custom models, or sophisticated multi-agent systems.

3.3 The Copilot family

"Copilot" spans several distinct products. Leaders routinely confuse them, so be precise:

Product	What it is	When you reach for it
Microsoft 365 Copilot	AI assistant inside Word, Excel, Outlook, Teams, PowerPoint; includes Chat, Agent Mode, and Cowork	Broad employee productivity; licensed per user
Copilot Chat	The conversational entry point; can tap enterprise data (e.g. via Fabric IQ)	Ad-hoc Q&A and grounded answers
Copilot Studio	Low-code platform to build & orchestrate custom agents and workflows	A bounded, repeatable task done by 50+ people
Copilot agents	The task-specific assistants you build in Copilot Studio	HR Q&A, ticket triage, support drafting, etc.
Agent 365	The governance + security control plane for agents at scale	Once you run several agents in production

2026 direction: Copilot Studio has shifted from "build a bot" to "operate a workforce of governed software actors" — with multi-agent orchestration, computer-using agents, real-time voice, and an administrable "agent estate." Model choice now includes frontier OpenAI and Anthropic models.

3.4 Microsoft Fabric, OneLake & Fabric IQ

AI is only as good as the data under it. Microsoft Fabric unifies data engineering, warehousing, real-time intelligence and Power BI on a single logical lake (OneLake) with zero-copy access.

- **Copilot in Fabric** helps author queries, pipelines and reports in natural language, respecting permission boundaries.
- **Fabric IQ** adds an ontology/semantic layer — a consistent business vocabulary on top of OneLake — so agents and copilots answer with shared, governed definitions. It is extending into Microsoft 365 Copilot (Chat and Cowork).
- **Data agents** let business users ask questions of governed data in plain English; operations agents can monitor metrics and act.

WATCH OUT

Two cost/quality traps to flag. (1) Copilot in Fabric consumes capacity units — heavy use during peak hours can cause throttling; consider isolated Copilot capacity. (2) Copilot for Power BI is only as good as your semantic model — weak models give weak answers. Data foundations come before clever agents.

3.5 Buy, configure, or build?

The most useful decision a leader makes here is which build mode fits the use case.

If the use case is…	Choose	Why
General productivity for many employees	Buy: Microsoft 365 Copilot	Fastest value, no build, per-user licensing
A bounded, repeatable task with defined data	Configure: Copilot Studio agent	Low-code, fast, governed via Agent 365
Custom models, complex multi-agent, pro-code	Build: Microsoft Foundry	Full control, evaluation, scale
Natural-language analytics on governed data	Fabric data agents + Fabric IQ	Grounded in your semantic layer

Hands-on · Stack-mapping + tool tour

HANDS-ON LAB

Goal. Map your Module 2 use case onto the Microsoft stack and (optionally) get hands on a real tool. ~30 minutes.

Part A · Stack map (everyone, ~15 min). On worksheet B4, for your top use case fill in: which APPLY/BUILD product; which models; what DATA in Fabric/OneLake; what retrieval (Foundry IQ / Azure AI Search); and which GOVERNANCE controls (Entra, Purview, Agent 365). Decide buy / configure / build.

Part B · Tool tour (if laptops + internet, ~15 min). Pick ONE:

- **Copilot Chat / Agent Mode:** prompt it to draft an artefact from your use-case brief; observe grounding and limits.
- **Copilot Studio (trial):** create a simple agent — give it instructions, point it at a knowledge source, test it in the playground. No code.
- **Microsoft Foundry portal:** browse the model catalog and the evaluations view to see what "governed build" looks like.

FACILITATOR NOTE

Setup & prerequisites. Part A needs nothing but the worksheet — always runnable. For Part B, arrange access in advance: a Microsoft 365 Copilot test tenant, a Copilot Studio trial, and/or a Foundry portal login. If access can't be guaranteed, run Part B as a projected live demo from the front instead of per-table.

Answer guidance for Part A. For the worked support example (Module 2): APPLY = a Copilot Studio agent surfaced in the agent's tooling (or M365 Copilot); MODELS = a frontier chat model from the Foundry catalog; DATA = knowledge base + tickets in OneLake; RETRIEVAL = Foundry IQ / Azure AI Search; GOVERNANCE = Entra for identity, Purview for data, human-in-the-loop. Build mode = configure (Copilot Studio). Reinforce: a quick win rarely needs pro-code Foundry.

USE CASE

Three stack-mapped examples.

- **HR policy Q&A:** Copilot Studio agent → HR docs in SharePoint/OneLake → Foundry IQ retrieval → Entra-scoped to employees.
- **Finance variance analysis:** Fabric data agent + Fabric IQ ontology over the finance semantic model → surfaced in Copilot Chat.
- **Multi-step claims handling:** Foundry multi-agent workflow → multiple tools/APIs → Agent 365 governance → human approval gate.

Knowledge check

1. Name the five layers of the Microsoft AI stack and what runs in each.
2. What changed about “Azure AI Foundry” in 2026, and what stayed the same?
3. For your use case, did you choose buy, configure, or build — and why?
4. Why are data foundations (Fabric / semantic models) a prerequisite, not an afterthought?

KEY TAKEAWAYS

- Five layers — Apply, Build, Data, Models, Infra — wrapped in governance. Know which box solves which problem.
- Microsoft Foundry (was Azure AI Foundry) = pro-code build; Copilot Studio = low-code; M365 Copilot = buy and use.
- Fabric + Fabric IQ give agents governed, shared data definitions — fix the data layer first.
- Default to the lightest build mode that fits: buy → configure → build.

Module 4 · ROI, Cost Models and Value Realization

Duration: 80 minutes · **Format:** Teach + build-a-business-case lab

Learning objectives

1. Explain why AI ROI is uniquely hard and how value realization actually works.
2. Build a complete cost stack — including the hidden and recurring costs people forget.
3. Quantify benefits across cost, productivity, revenue, risk, and experience.
4. Assemble a defensible business case and define metrics that prove (or disprove) value.

FACILITATOR NOTE

Timing. (1) Why ROI is hard – 8 min. (2) The cost stack – 12 min. (3) Benefit categories – 12 min. (4) Value-realization loop & metrics – 8 min. (5) Lab: build the model – 25 min. (6) Debrief + check – 10 min.

Tone. Leaders have heard inflated ROI claims. Be the calm, numerate voice: ranges not hero-numbers, total cost not just licence cost, and an honest “we’ll measure and kill it if it doesn’t work.” That credibility is what unlocks funding.

4.1 Why AI ROI is hard

- **The pilot trap.** PoCs are cheap and impressive; production, integration, change management and run-costs are where the money and the value both live. Most initiatives that disappoint stalled between pilot and production.
- **Diffuse benefits.** “30 minutes saved per person per day” is real but evaporates unless the freed capacity is redeployed to something that shows up in a P&L.
- **Moving costs.** Token/usage pricing means cost scales with success — a popular agent can blow its budget. ROI is a curve, not a point.

This is why a meaningful share of AI/agentic projects get cancelled — not because the tech failed, but because value was never defined, measured, or captured.

4.2 The cost stack

A credible business case prices the whole stack, not just the licence. Costs fall into build (one-off) and run (recurring).

Cost layer	Examples	Build / Run
Platform & licences	M365 Copilot seats, Copilot Studio, Fabric capacity, Azure	Run
Model / token usage	Inference per token; scales with adoption	Run (variable)
Data & integration	Pipelines, connectors, semantic models, cleanup	Build + Run
Build effort	Design, prompt/agent engineering, evaluation, testing	Build
People & change	Training, adoption, redesigned roles & processes	Build + Run

Cost layer	Examples	Build / Run
Run & operate	Monitoring, drift, support, FinOps, re-evaluation	Run
Governance & risk	Reviews, controls, audit, compliance effort	Build + Run

WATCH OUT

The three most-forgotten costs: (1) change management — often larger than the build; (2) variable token/usage cost at scale; (3) ongoing evaluation & governance. Omit these and your ROI is fiction.

4.3 Benefit categories

Map every claimed benefit to a category and a number with a baseline. "Better" is not a benefit; "first-response time from 9h to 2h" is.

Category	Typical metric	How to capture the value
Cost avoidance	Hours saved × loaded rate	Only counts if capacity is redeployed or not hired
Productivity	Throughput, cycle time	Tie to output that matters, not activity
Revenue / growth	Conversion, retention, deal size	Attribute carefully; use a control group
Quality / risk	Error rate, rework, incidents	Fewer errors = direct + reputational value
Experience	CSAT, eNPS, response time	Leading indicator of retention & growth

4.4 The value-realization loop

Value is not realized at go-live; it is realized by measuring, learning, and reinvesting. Spend becomes capability, capability produces outcomes, outcomes become value — and the loop redirects money to what works.

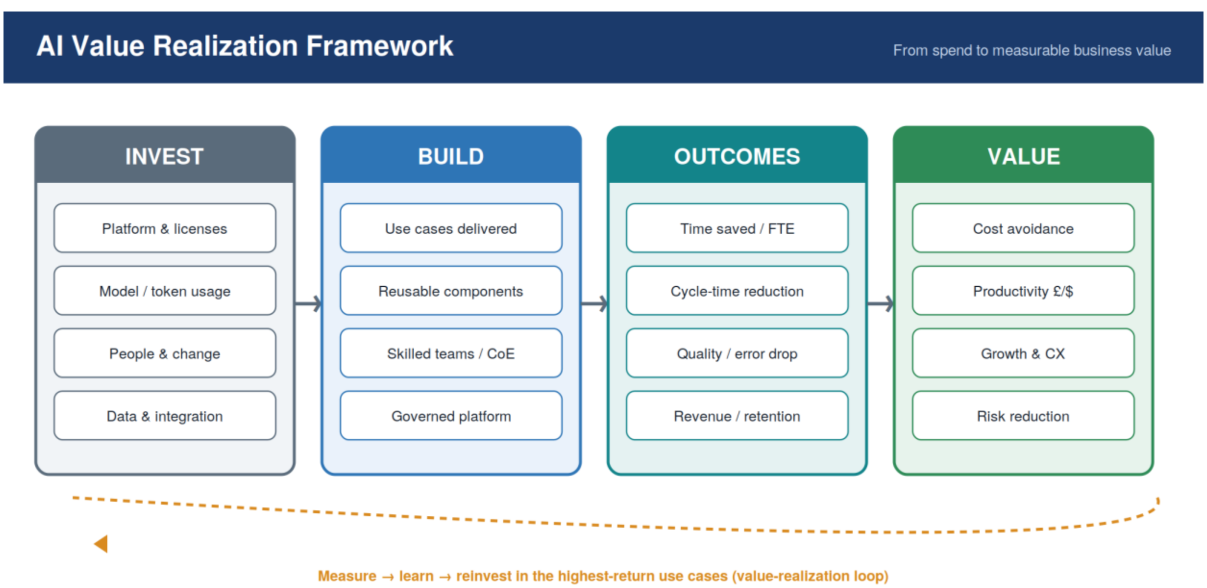

Figure 4.1 — AI value-realization framework.

4.5 A worked business case

Using the support use case (9h → target 2h first-response), a simple year-one model for a 40-agent team:

Line	Assumption	Annual figure
Time saved	40 agents × ~1.0h/day × 220 days	8,800 hours
Value of time	8,800h × $35 loaded rate (redeployed)	$308,000
Retention uplift	Lower handle time → modest churn reduction (modelled)	$60,000
Gross benefit	—	**$368,000**
Build cost	Agent build + data + evaluation (one-off)	$90,000
Run cost	Licences + tokens + monitoring + governance	$70,000
Net year-1 value	—	**≈ $208,000**

Illustrative only. The point is the structure — benefits with baselines, the full cost stack, and a redeployment assumption stated honestly — not these specific numbers.

4.6 Metrics that prove value

- **Define the ONE primary metric** and its baseline before you build. If it can't be measured, the use case isn't ready.
- **Pair leading and lagging indicators:** adoption and quality (leading) predict the cost/revenue movement (lagging).
- **Track cost per outcome** (e.g. $ per resolved ticket), not just total spend — this is your FinOps early-warning.

WATCH OUT

Avoid vanity metrics. "Number of prompts," "users onboarded," or "messages sent" measure activity, not value. A use case can be wildly "engaged" and still move nothing that matters.

Hands-on · Build the business case

HANDS-ON LAB

Goal. Produce a one-page ROI model for your table's top use case. ~25 minutes. Use worksheet B5.

1. State the primary metric and its baseline.
2. Quantify benefits by category — be conservative; state every assumption (especially redeployment of saved time).
3. Price the full cost stack: build (one-off) and run (recurring, including variable token cost).
4. Compute net year-1 value and a rough payback period.
5. Write the one-sentence value narrative a CFO would accept.

FACILITATOR NOTE

Facilitation. Hand each table the worked example (4.5) as a template. Pressure-test two things: (a) “where does the saved time actually go?” — if nobody can answer, the cost-avoidance benefit is soft; (b) “what’s your token/usage cost at 10× adoption?” — surfaces the variable-cost trap. A model that survives both questions is fundable.

USE CASE

Industry ROI signals (cite as ranges, verify before quoting).

- Microsoft 365 Copilot: independent analysis has reported triple-digit three-year ROI and several hours saved per user per month — contingent on real adoption and redeployment.
- Generative AI overall: IDC/Microsoft have cited an average ~3.7x return per $1 invested — averages hide wide variance.
- Counterweight: an IBM CEO study found only ~25% of AI initiatives met ROI expectations — the difference is disciplined value capture.

Knowledge check

1. What are the three most commonly forgotten costs in an AI business case?
2. For your use case, what is the ONE primary metric and its baseline?
3. Why can token/usage pricing make ROI a curve rather than a fixed number?
4. Give a vanity metric for your use case — and the value metric you’ll use instead.

KEY TAKEAWAYS

- ROI is hard because value leaks between pilot and production — plan for production and run-costs.
- Price the whole cost stack: build + run, including change management, variable tokens, and governance.
- Every benefit needs a category, a number, and a baseline — and an honest redeployment story.
- Value is realized through a measure–learn–reinvest loop; track cost per outcome, not vanity metrics.

Module 5 · AI Governance and Responsible AI Principles

Duration: 80 minutes · **Format:** Teach + risk-tiering activity

Learning objectives

1. Articulate why governance is now a board-level concern, not a compliance afterthought.
2. Apply the Responsible-AI principles to a real use case.
3. Set up a right-sized governance operating model with clear roles and decision gates.
4. Risk-tier use cases and identify the controls and Microsoft tooling each tier needs.

FACILITATOR NOTE

Timing. (1) Why governance now – 10 min. (2) Responsible-AI principles – 12 min. (3) Operating model & gates – 13 min. (4) Microsoft governance tooling – 8 min. (5) Agentic-specific risks – 7 min. (6) Risk-tiering activity – 20 min. (7) Debrief + check – 5 min.

Framing. Position governance as the enabler of speed, not the brake. "Light enough not to block, real enough to be trusted" is the line that lands with leaders who fear bureaucracy.

5.1 Why governance now

- **Autonomy raises the stakes.** Agents that act — not just answer — can cause real-world harm: wrong transactions, data exposure, biased decisions at scale.
- **A new attack surface.** Tool access, broad permissions, and memory create risks (e.g. prompt injection, data exfiltration) most security teams aren't yet equipped for; agent-linked incidents are now being tracked.
- **Regulation is arriving.** Risk-based regimes — the EU AI Act (phased obligations through 2025–2027), plus frameworks such as the NIST AI Risk Management Framework and ISO/IEC 42001 — expect documented risk management, transparency, and human oversight. Confirm the current status and your jurisdiction's rules before advising.
- **Trust is the growth constraint.** Customers, employees and boards adopt AI only as fast as they trust it. Governance is how you earn the licence to scale.

5.2 Responsible-AI principles

Microsoft's six principles are a practical checklist for any use case. Accountability and transparency sit underneath the other four — someone is answerable, and the system is explainable.

Figure 5.1 — The six Responsible-AI principles.

Principle	Ask of every use case
Fairness	Could this disadvantage a group? How will we test for bias?
Reliability & safety	How does it behave on bad input or at the edges? What's the fail-safe?
Privacy & security	What data does it touch? Who can access it? Is it minimized and protected?
Inclusiveness	Does it work for everyone who must use it, including accessibility needs?
Transparency	Do users know it's AI? Can we explain a given output or decision?
Accountability	Who is the named owner answerable for this system in production?

5.3 Governance operating model

Right-sized governance assigns clear roles and puts decision gates on the use-case lifecycle — so speed and safety are not in tension.

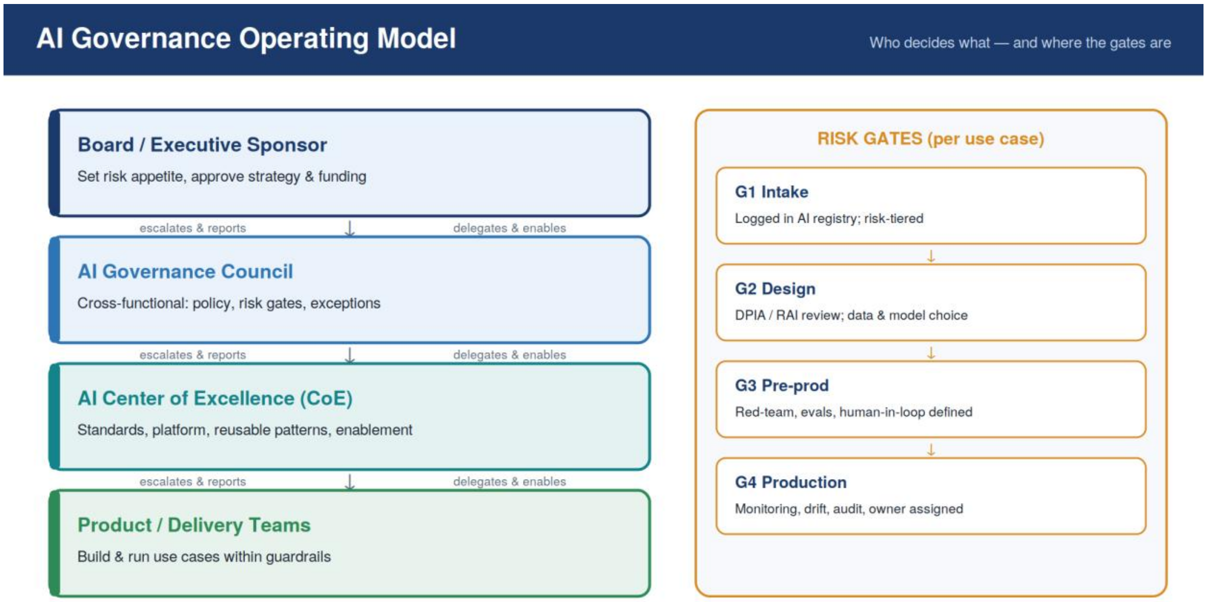

Figure 5.2 — AI governance operating model: roles (left) and lifecycle risk gates (right).

- **Board / executive sponsor** — sets risk appetite, approves strategy and funding.
- **AI governance council** — cross-functional (risk, legal, security, data, business); owns policy, gate decisions and exceptions.
- **AI Center of Excellence** — standards, the platform, reusable patterns, enablement.
- **Product / delivery teams** — build and run use cases within the guardrails.

Lifecycle gates: G1 Intake (registered & risk-tiered) → G2 Design (impact assessment, data/model choice) → G3 Pre-production (red-team, evaluations, human-in-the-loop defined) → G4 Production (monitoring, drift, audit, named owner).

5.4 Risk tiering

Not every use case needs the same scrutiny. Tier by potential impact, then apply proportionate controls. Over-governing a low-risk copilot wastes goodwill; under-governing an autonomous agent invites disaster.

Tier	Examples	Controls
Low	Internal drafting, summarization, search assist	Usage policy, basic logging, light review
Medium	Customer-facing assistant, internal decision support	Human-in-the-loop, evaluations, monitoring, DPIA
High	Autonomous action, regulated decisions, vulnerable users	Council approval, red-team, strict oversight, audit, kill-switch

5.5 Microsoft tooling for governance

Governance is operationalized with the same stack from Module 3:

- **Microsoft Entra** — identity for users and agents; scoping and conditional access.
- **Microsoft Purview** — data classification, DLP, compliance, and visibility into AI data use.
- **Agent 365** — the security/governance plane for agents at scale (discovery, posture, controls).
- **Foundry evaluations + Content Safety** — testable safety/quality checks and guardrails at build time and in production.
- **Defender** — threat protection across the AI surface.

5.6 Agentic-AI-specific risks

WATCH OUT

What's different once an agent can act:

- **Identity & permissions:** an agent acting on a user's behalf needs its own scoped identity and least-privilege access — not a shared admin token.
- **Tool access:** every tool/API an agent can call is a way to cause effects; allow-list and gate them.
- **Memory:** persistent memory can leak data across sessions or users if not isolated.
- **Autonomy bounds:** define what the agent may do unsupervised vs. what requires human approval; always have a kill-switch.
- **Multi-agent emergence:** agents calling agents can produce behaviour no single component intended — test the system, not just the parts.

Hands-on · Risk-tier and gate your use case

HANDS-ON LAB

Goal. Assign a risk tier and design the gates for your table's top use case. ~20 minutes. Use worksheet B6.

1. Score impact: data sensitivity, autonomy level, who is affected, regulatory exposure, reversibility of a wrong action.
2. Assign a tier (low / medium / high) from the highest-scoring factor.
3. Run the six Responsible-AI questions (5.2); note any that raise a flag.
4. Define each gate (G1–G4): what must be true to pass, and who signs off.
5. Name the accountable owner and the kill-switch.

FACILITATOR NOTE

Answer guidance — the support use case. Likely **medium** tier: customer-facing and uses customer data, but a human sends every reply (no autonomous action) and the decision is reversible. Controls: human-in-the-loop, Purview on the data, evaluations for accuracy/tone, monitoring, a DPIA. It would jump to **high** the moment you let it auto-send or take account actions — a great live illustration of how autonomy drives tier.

Facilitation. Ask each table: “what one change would push this up a tier?” It makes the autonomy–risk link visceral and ties the whole day together.

USE CASE

Governance scenarios to discuss if time allows:

- **The eager builder:** a team ships an unregistered agent with broad data access. Which gate failed, and what’s the fix? (G1 intake; require registration + tiering before build.)
- **The helpful over-sharer:** a copilot surfaces salary data to the wrong employee. Which principle and control? (Privacy & security; Purview + Entra scoping.)
- **The autonomous escalation:** a procurement agent places an erroneous order. What was missing? (Autonomy bounds, human approval gate, kill-switch.)

Knowledge check

1. Why is governance an enabler of speed rather than a brake on it?
2. Walk one use case through the six Responsible-AI questions — which raises the biggest flag?
3. What risk tier is your use case, and what single change would push it up a tier?
4. Name the four lifecycle gates and one thing that must be true to pass each.

KEY TAKEAWAYS

- Governance is now board-level: autonomy, a new attack surface, regulation, and trust all demand it.
- Six Responsible-AI principles — with accountability and transparency as the foundation — are your per-use-case checklist.
- Right-size it: clear roles + four lifecycle gates + risk tiering, operationalized with Entra, Purview, Agent 365.
- Autonomy drives risk tier. The moment an agent can act, the controls step up.

Day 1 · Wrap-up and bridge to Day 2

What you built today

Each table now holds a single, continuous thread — one real business problem that has been:

1. Placed honestly on the maturity curve (Module 1).
2. Identified, qualified, and prioritized as a use case (Module 2).
3. Mapped onto the Microsoft AI stack with a buy/configure/build decision (Module 3).
4. Costed and turned into a defensible business case (Module 4).
5. Risk-tiered with governance gates and an accountable owner (Module 5).

That end-to-end artefact — from problem to governed, funded plan — is the real output of Day 1.

Action items for sponsors

#	Action	Owner	By
1	Confirm the executive sponsor and the use-case business owner	Leadership	This week
2	Validate the top use case's data access and quality	Data owner	2 weeks
3	Stand up (or convene) a lightweight AI governance council	Sponsor	2 weeks
4	Pressure-test the business case with Finance	Owner + Finance	3 weeks
5	Decide go / no-go for a time-boxed production pilot	Council	1 month

Looking ahead

Day 1 aligned strategy and leadership. Subsequent days typically go deeper into hands-on building, data foundations, agent engineering, security, and operating model — turning today's prioritized, governed use case into a working pilot. Keep your one-page brief, ROI model, and risk tiering; they are the inputs to everything that follows.

FACILITATOR NOTE

Close strong. Do a round-the-room one-word check-out ("one word for how you feel about your team's use case"). Capture the parking-lot items and assign owners. Remind everyone the deliverable travels with them to Day 2 — it is not a throwaway exercise.

Appendix A · Facilitator checklist

Before the session

- Confirm audience seniority and tailor the depth (leadership = conceptual, not SDKs).
- Print worksheets B1–B6, one set per participant; bring spares.
- Sticky notes, markers, flip-charts; draw a 2×2 grid on a wall for Module 2.
- Arrange Module 3 tool access in advance (M365 Copilot tenant / Copilot Studio trial / Foundry portal); have a recorded demo as fallback.
- Load the eight figures in this guide for projection.
- Re-verify Microsoft product names, GA/preview status, and any pricing on Microsoft Learn — these change monthly.
- Ask each table to come with one real business problem (the golden thread).

During the session

- Protect the workshop time in Module 2 — it produces the artefact for the rest of the day.
- Keep a visible parking lot; don't let one expert pull the room into the weeds.
- Use ranges, not hero-numbers, for every statistic; attribute the source.
- Tie every module back to each table's single use case.
- Watch energy after lunch — Module 4 is numerate; keep it brisk and concrete.

After the session

- Photograph every wall (matrices, briefs) and share back within 24 hours.
- Circulate the action-item table with named owners and dates.
- Send the consolidated use-case briefs, ROI models, and risk tiers to sponsors.

Appendix B · Participant worksheets

Reproduce one set per participant. Each maps to a hands-on activity.

B1 · Maturity self-assessment (Module 1)

Pillar	Score 1–5	Evidence / note
Strategy & sponsorship		
Data readiness		
People & skills		
Governance		
Production footprint		
Honest stage = lowest score above		

Our stage is ___, and our single next move is ________________________________.

B2 · One-page use-case brief (Module 2)

Field	Your answer
Title & owner	
Problem (with a number)	
Today's process / friction	
AI pattern (predict/gen/agentic) + why	
Data needed (source/owner/quality/sensitivity)	
Success metric + baseline	
Value (1–5) / Feasibility (1–5) / quadrant	
Top risks / compliance	

B3 · Prioritization scorecard (Module 2)

Use case	Value 1–5	Feasibility 1–5	Data readiness	Quadrant

B4 · Stack map (Module 3)

Layer	Your choice for this use case
Apply (M365 Copilot / Copilot agent)	
Build (Copilot Studio / Foundry)	
Models (catalog / frontier / fine-tune)	
Data (Fabric / OneLake / Fabric IQ)	
Retrieval (Foundry IQ / Azure AI Search)	
Governance (Entra / Purview / Agent 365)	
Build mode: buy / configure / build	

B5 · ROI model (Module 4)

Line	Assumption	Annual figure
Primary metric + baseline		
Benefit 1 (category)		
Benefit 2 (category)		
Gross benefit		
Build cost (one-off)		
Run cost (recurring incl. tokens)		
Net year-1 value / payback		

B6 · Risk tiering & gates (Module 5)

Factor	Assessment
Data sensitivity	
Autonomy level (assist / decide / act)	
Who is affected	
Regulatory exposure	
Reversibility of a wrong action	
Assigned tier (low/med/high)	
RAI flags (from the six questions)	
Accountable owner + kill-switch	

Appendix C · Knowledge-check answer key

These checks are reflective, not pass/fail — most answers are organization-specific. Look for the reasoning below.

Module 1

1. Adoption uses AI within an unchanged process; transformation redesigns the process/business model.

2. Predictive = forecast/score; generative = draft/summarize; agentic = multi-step act. Any valid mapping.

3. Stage = lowest pillar; blocker is usually data, governance, or getting one thing into production.

4. More agents add cost, fragility, and governance load; use the simplest pattern — a rule or single model call often wins.

Module 2

1. Value chain, pain points, data-rich processes, decisions-at-scale, scarce knowledge.

2. Deterministic logic, zero-tolerance-without-oversight, missing data, or trivial value → rules/automation/BI.

3. Separating value and feasibility prevents one strong dimension masking a fatal weakness and forces a real trade-off.

4. Organization-specific; insist on a real baseline number.

Module 3

1. Apply, Build, Data, Models, Infra — wrapped in governance; recall the example products in each.

2. Azure AI Foundry was renamed Microsoft Foundry (Jan 2026); capabilities unchanged.

3. Buy for broad productivity, configure for bounded tasks, build for custom/complex — justify against the use case.

4. Agents answer from data; weak/ungoverned data → wrong or unsafe answers, so data comes first.

Module 4

1. Change management, variable token/usage cost at scale, ongoing evaluation & governance.

2. Organization-specific; must include a baseline.

3. Usage-based pricing makes cost scale with adoption, so ROI moves along a curve.

4. Vanity = activity (prompts, logins); value = outcome moved (cost/time/revenue/quality).

Module 5

1. Clear, proportionate gates let teams move fast safely and build the trust needed to scale.

2. Any principle is acceptable; look for a concrete control attached to the flag.

3. Tier from the highest-impact factor; autonomy or regulated data usually pushes it up.

4. G1 Intake (registered/tiered), G2 Design (impact assessment), G3 Pre-prod (red-team/evals/HITL), G4 Production (monitoring/owner).

Appendix D · Glossary

Term	Meaning
Agent	An AI system that plans and takes multi-step actions, often calling tools or APIs.
Agent 365	Microsoft's governance/security control plane for operating agents at enterprise scale.
Agentic AI	AI that acts autonomously across steps — the third wave after predictive and generative.
Copilot Studio	Microsoft's low-code platform for building and orchestrating custom agents.
CoE	Center of Excellence — the team owning AI standards, platform, and reusable patterns.
DPIA	Data Protection Impact Assessment — a structured privacy-risk review.
Fabric IQ	The semantic/ontology layer in Microsoft Fabric giving agents shared business definitions.
Foundry IQ	Microsoft Foundry's SLA-backed retrieval/knowledge plane for grounding agents.
Frontier model	A leading-edge large model (e.g. top OpenAI or Anthropic models) in the catalog.
Generative AI	AI that creates content — text, code, images — or summarizes and answers.
Human-in-the-loop	A required human review/approval step before an AI output takes effect.
Microsoft Foundry	Microsoft's unified AI build platform (renamed from Azure AI Foundry, Jan 2026).
MCP	Model Context Protocol — a standard for connecting agents to tools and data sources.
OneLake	Microsoft Fabric's single logical data lake with zero-copy access.
Predictive AI	AI that forecasts or scores from historical patterns — the first wave.
Purview	Microsoft's data governance, classification, and compliance platform.
RAG	Retrieval-Augmented Generation — grounding a model's answers in retrieved enterprise data.
Responsible AI	Microsoft's six-principle framework: fairness, reliability & safety, privacy & security, inclusiveness, transparency, accountability.
Risk tier	A low/medium/high classification that sets how much governance a use case needs.
Value realization	Turning AI spend into measured business value via a measure–learn–reinvest loop.

Appendix E · References & further reading

Primary, regularly-updated sources. Verify currency before quoting figures or product details.

- Microsoft Learn — Microsoft Foundry, Copilot Studio, Microsoft Fabric & Fabric IQ documentation (product capabilities, GA/preview status).
- Microsoft Foundry Blog & Microsoft Copilot Blog — monthly “What’s new” updates (e.g. Build 2026 announcements).
- Microsoft Responsible AI Standard — the six principles and implementation guidance.
- Forrester / IDC — Microsoft 365 Copilot and generative-AI ROI analyses (cite as ranges).
- McKinsey Global AI Survey 2026 — adoption, ROI, and high-performer practices.
- Gartner — forecasts on enterprise agent embedding, project cancellation rates, and agent ecosystems.
- IDC Worldwide AI Spending Guide — market size and spend growth.
- NIST AI Risk Management Framework and ISO/IEC 42001 — governance frameworks.
- EU AI Act — official text and phased-obligation timeline (confirm current applicability for your jurisdiction).

Currency note. *This guide reflects the landscape as of June 2026 (Microsoft Build 2026). The AI market and Microsoft’s product names, availability, and pricing change frequently — always re-verify against Microsoft Learn and primary analyst sources before presenting to a client.*